

Value-based Reinforcement Learning

Some Discussions

Kan Ren

Apex Data and Knowledge Management Lab
Shanghai Jiao Tong University

Aug. 3 2017



Outline

1 Revision of Value-based RL

- Sarsa & Q-learning
- Q-learning

2 Issues in Q-learning

- Overestimation
- Double Q-learning
- Averaged Q-learning



Outline

1 Revision of Value-based RL

- Sarsa & Q-learning
- Q-learning

2 Issues in Q-learning

- Overestimation
- Double Q-learning
- Averaged Q-learning



Sarsa

Algorithm

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal



Outline

1 Revision of Value-based RL

- Sarsa & Q-learning
- Q-learning

2 Issues in Q-learning

- Overestimation
- Double Q-learning
- Averaged Q-learning



Q-Learning

Algorithm

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal



Difference

Update Rule

- Sarsa



Outline

1 Revision of Value-based RL

- Sarsa & Q-learning
- Q-learning

2 Issues in Q-learning

- Overestimation
- Double Q-learning
- Averaged Q-learning



Outline

1 Revision of Value-based RL

- Sarsa & Q-learning
- Q-learning

2 Issues in Q-learning

- Overestimation
- Double Q-learning
- Averaged Q-learning



Outline

1 Revision of Value-based RL

- Sarsa & Q-learning
- Q-learning

2 Issues in Q-learning

- Overestimation
- Double Q-learning
- Averaged Q-learning

